

Data Science: Tidyverse

Alex Di Genova

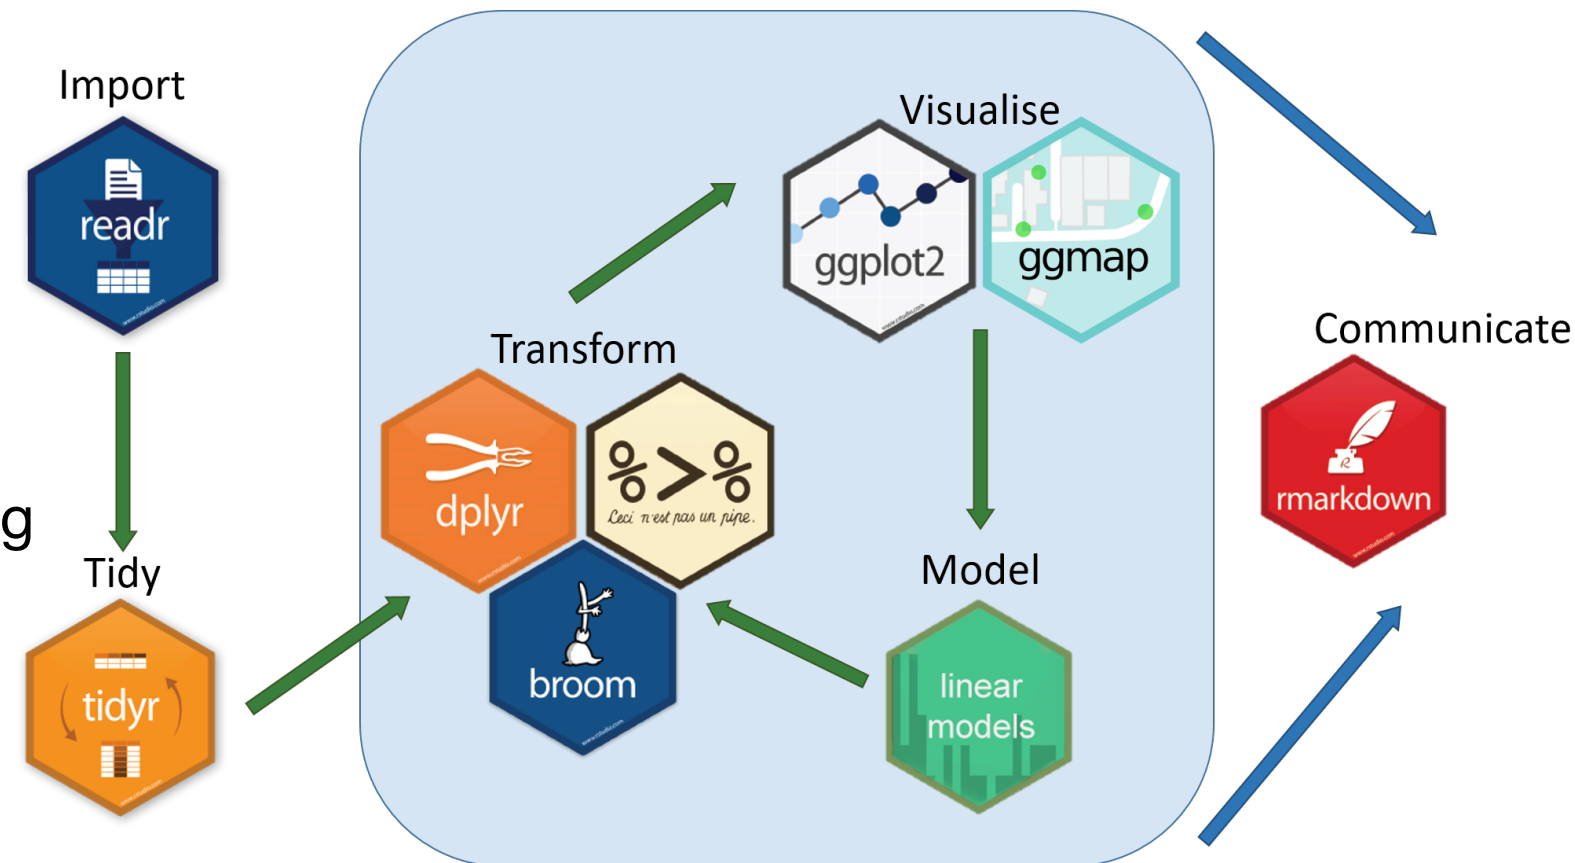
16/05/2024

What is Tidyverse?

- A collection of R packages for data science

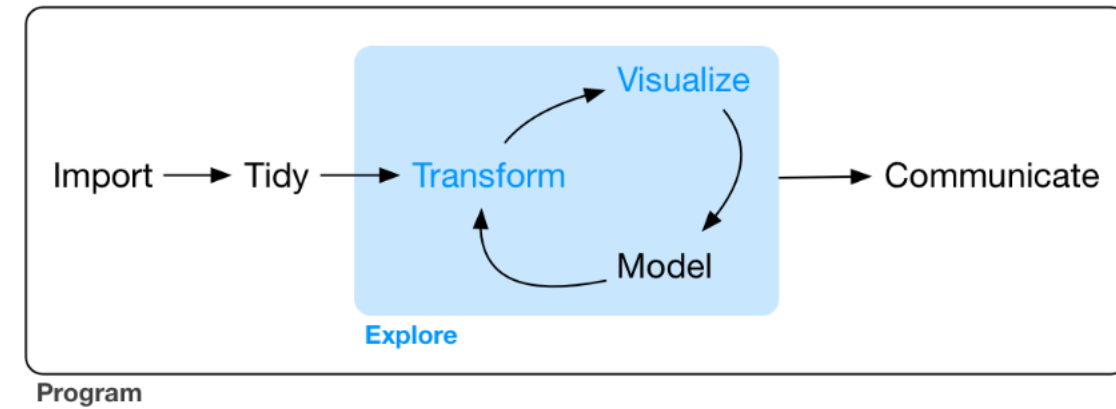
- ggplot2 — data visualization
- dplyr — data manipulation
- tidyr — data tidying
- readr — data import
- purrr — functional programming
- tibble — modern dataframes
- stringr — string manipulation
- forcats — factor handling

- Data science workflow (import, clean, transform, visualize, model)



Tidyverse

Explore




5. Place geoms in a cartesian coordinate system.

6. Map the y values to **..count..** and the x values to **cut**.

carat	cut	color	clarity	depth	table	price	x	y	z
0.23	Ideal	E	SI2	61.5	55	326	3.95	3.98	2.43
0.21	Premium	E	SI1	59.8	61	326	3.89	3.84	2.31
0.23	Good	E	VS1	56.9	65	327	4.05	4.07	2.31
0.29	Premium	I	VS2	62.4	58	334	4.20	4.23	2.63
0.31	Good	J	SI2	63.3	58	335	4.34	4.35	2.75
...

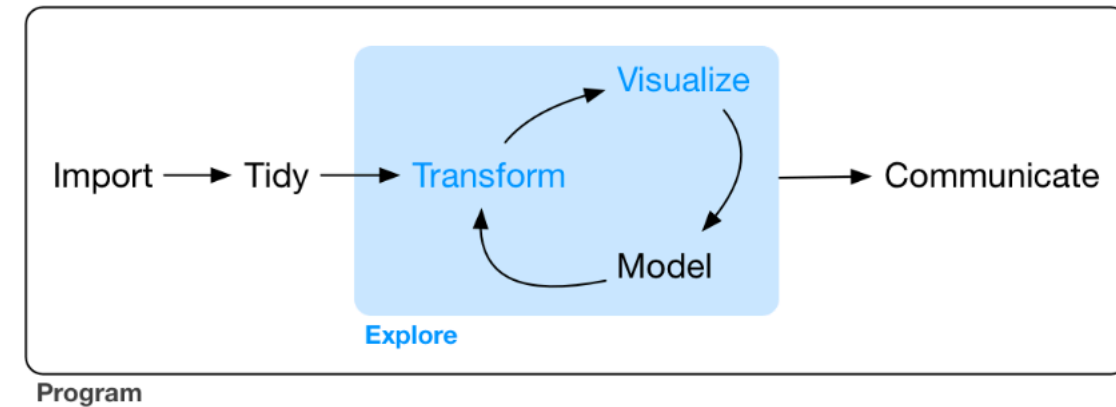
stat_count()

cut	count	prop
Fair	1610	1
Good	4906	1
Very Good	12082	1
Premium	13791	1
Ideal	21551	1



Tidyverse


Dplyr — transform



- Data manipulation challenges:
 - Pick observations by their values (`filter()`)
 - Reorder the rows (`arrange()`)
 - Pick variables by their names (`select()`)
 - Create new variables (`mutate()`)
 - Collapse many values to a single summary (`summarize()`)
 - Group rows (`group_by()`)
- All dplyr verbs expect a data.frame and produce a new data.frame

Dplyr — transform


Select



- Select columns with `select()`
 - `mpg %>% select(model,manufacturer,cyl)`
 - `mpg %>% select(model:year)`
 - `mpg %>% select(-(model:year))`
- Helper functions
 - `starts_with("abc")` matches names that begin with “abc”.
 - `ends_with("xyz")` matches names that end with “xyz”.
 - `contains("ijk")` matches names that contain “ijk”.
 - `num_range("x", 1:3)` matches x1, x2, and x3
- `rename(modelo = model)`
- `mpg %>% select(trans, drv, everything())`

Dplyr — transform

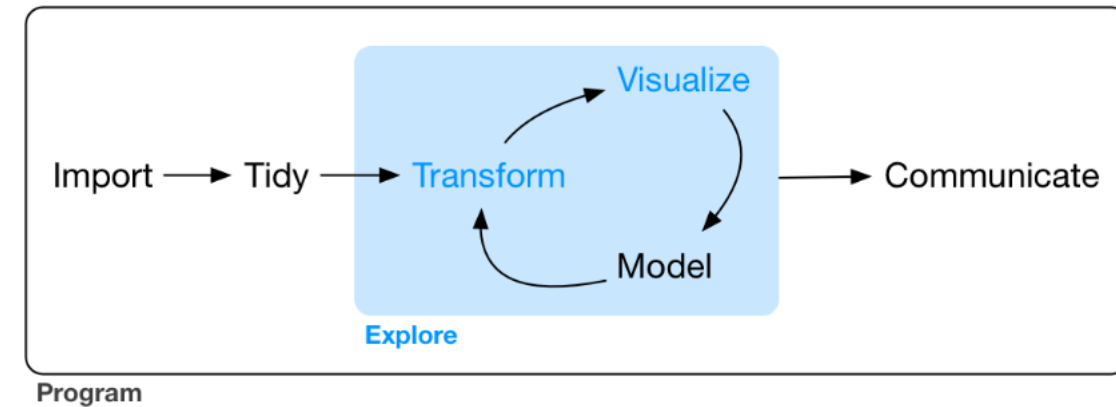
Mutate



- New Variables with `mutate()`
- `mutate()` always adds new columns at the end of data frame
 - `mpg %>% mutate(delta=hwy-cty)`
 - `?transmute()`
- Helper functions
 - Arithmetic operators `+`, `-`, `*`, `/`, `^`
 - Logs `log()`, `log2()`, `log10()`
 - Cumulative R provides `cumsum()`, `cumprod()`, `cummin()`, `cummax()`; and dplyr provides `cummean()`
 - Logical comparisons `<`, `<=`, `>`, `>=`, `!=`

Dplyr — transform


Summarise



- Grouped Summaries with `summarize()`
- collapses a data frame to a single row
 - `mpg %>% summarise(mean(hwy))`
- `summarize()` and `group_by()` a perfect couple
 - `mpg %>% group_by(manufacturer) %>% summarize(m=mean(hwy)) %>% arrange(desc(m))`
- Helper functions
 - Measures of location: `mean(x)`, `median(x)`
 - Measures of spread: `sd(x)`, `IQR(x)`, `mad(x)`
 - Measures of rank: `min(x)`, `quantile(x, 0.25)`, `max(x)`
 - Counts: `n()`, `n_distinct(x)`

Dplyr — transform

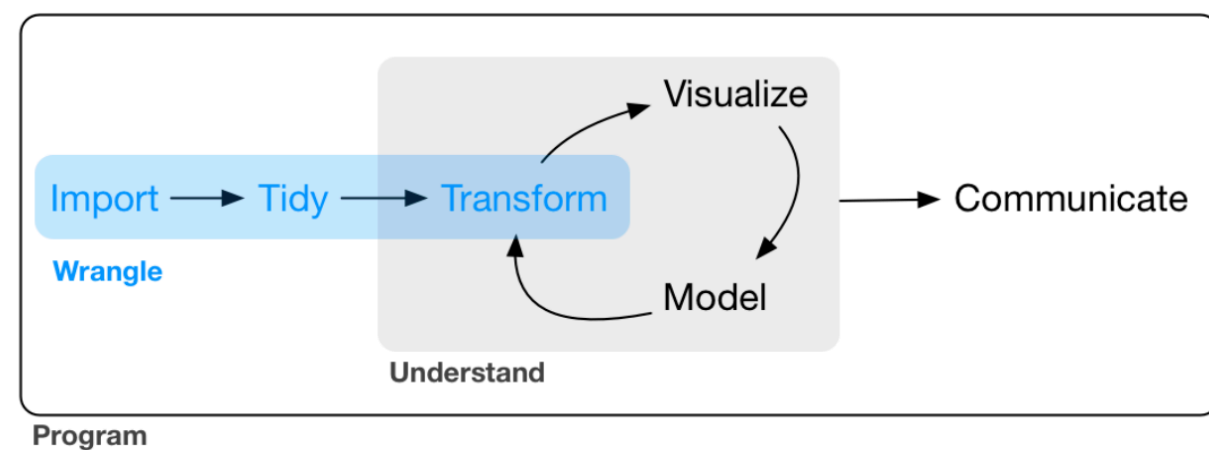
Summarise



- Grouped Summaries with `summarize()`
- collapses a data frame to a single row
 - `mpg %>% summarise(mean(hwy))`
- `summarize()` and `group_by()` a perfect couple
 - `mpg %>% group_by(manufacturer) %>% summarize(m=mean(hwy)) %>% arrange(desc(m))`
- Helper functions
 - Measures of location: `mean(x)`, `median(x)`
 - Measures of spread: `sd(x)`, `IQR(x)`, `mad(x)`
 - Measures of rank: `min(x)`, `quantile(x, 0.25)`, `max(x)`
 - Counts: `n()`, `n_distinct(x)`

Tidyverse

Wrangle



- Getting your data in a helpful form for visualization and modeling

- data.frames and Tibbles

- Data import with readr

```
### readr
```

```
`{r readdata}
```

```
d = read_csv("../data/worldcitiespop.csv.gz")
```

```
`{r}
```

Rows: 3173958 Columns: 7 — Column specification

Delimiter: ","

- From a file.txt to a data.frame

- Common functions

- read_csv(). read_delim(), read_delim()

Rows: 3,173,958

Columns: 7

\$ Country <chr> "ad", "ad", "ad", "ad", "ad", "ad", "

\$ City <chr> "aixas", "aixirivali", "aixirivall",

\$ AccentCity <chr> "Aixàs", "Aixirivali", "Aixirivall",

\$ Region <chr> "06", "06", "06", "06", "06", "07", "

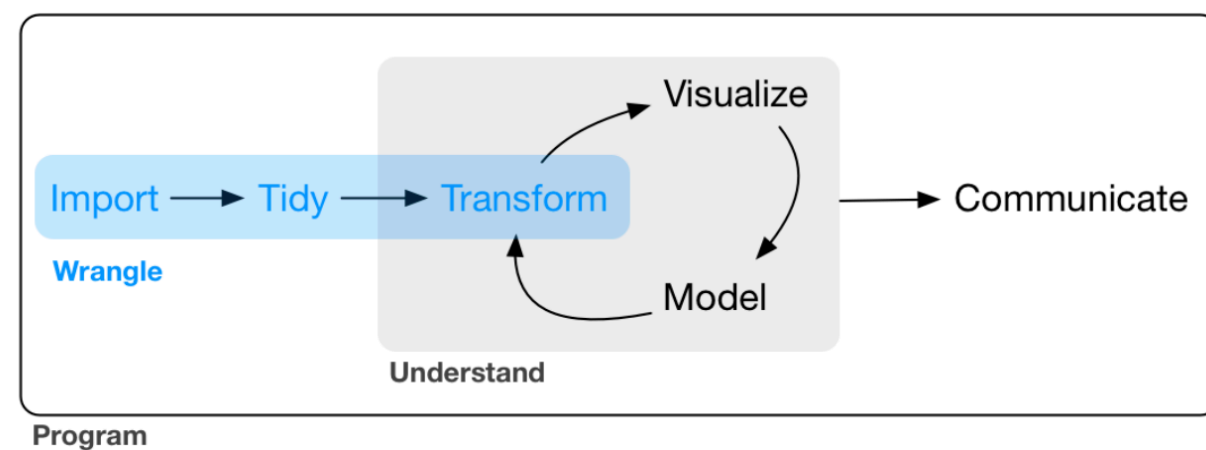
\$ Population <dbl> NA, NA, NA, NA, NA, NA, 20430, NA, NA

\$ Latitude <dbl> 42.48333, 42.46667, 42.46667, 42.4666

\$ Longitude <dbl> 1.466667, 1.500000, 1.500000, 1.500000

Tidyverse

Readr



```
read_csv("The first line of metadata
The second line of metadata
x,y,z
1,2,3", skip = 2)
```

```
read_csv("# A comment I want to skip
x,y,z
1,2,3", comment = "#")
```

```
read_csv("1,2,3\n4,5,6", col_names = c("x", "y", "z"))
```

- `parse_*` functions:

- From a character vector to a specialized vector like `date` or `double`.

- `parse_logical()`, `parse_date()`

```
```{r fechas}
parse_date("01/02/15", "%m/%d/%Y")
parse_date("01/02/15", "%d/%m/%Y")
parse_date("01/02/15", "%Y/%m/%d")
parse_date("12 Enero 2015", "%d %B %Y")
```

How readr automatically guesses

```
[1] "2015-01-02"
[1] "2015-02-01"
[1] "2001-02-15"
[1] "2015-01-12"
```

```
```{r n}
parse_number("$100")
parse_number("20%")
parse_number("It cost $123.45")
```
```

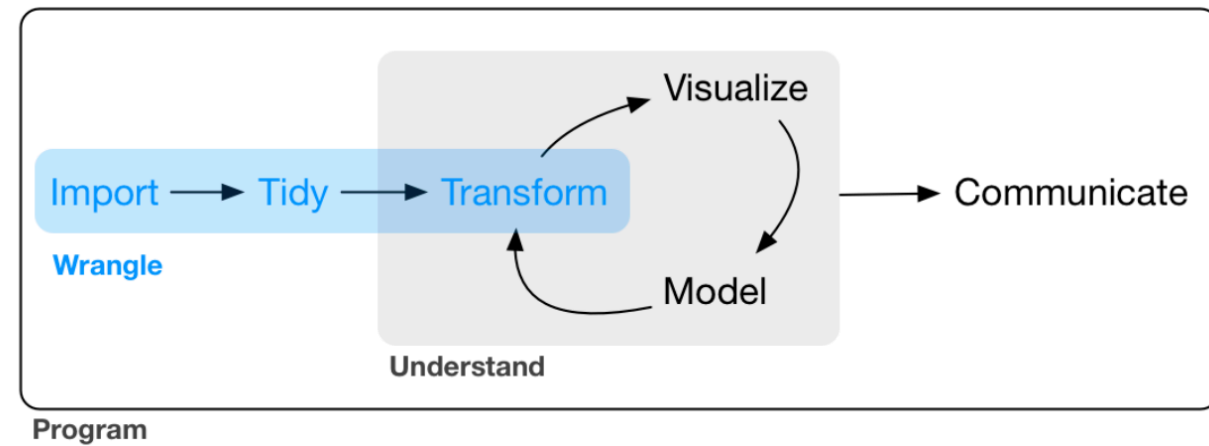
```
[1] 100
[1] 20
[1] 123.45
```

```
```{r guess}
guess_parser("2010-10-01")
guess_parser("15:01")
guess_parser(c("TRUE", "FALSE"))
guess_parser(c("1", "5", "9"))
guess_parser(c("12,352,561"))
```
```

```
[1] "date"
[1] "time"
[1] "logical"
[1] "double"
[1] "number"
```

# Tidyverse

## Readr



- `write_csv()` and `write_tsv()`
  - `write_csv(challenge, "challenge.csv")`
- `write_rds()` and `read_rds()`, store data in R's custom binary format called RDS
  - `write_rds(challenge, "challenge.rds")`
  - `read_rds("challenge.rds")`

# Tidyverse

## Tidy

*Journal of Statistical Software*

|              | treatmenta | treatmentb |
|--------------|------------|------------|
| John Smith   | —          | 2          |
| Jane Doe     | 16         | 11         |
| Mary Johnson | 3          | 1          |

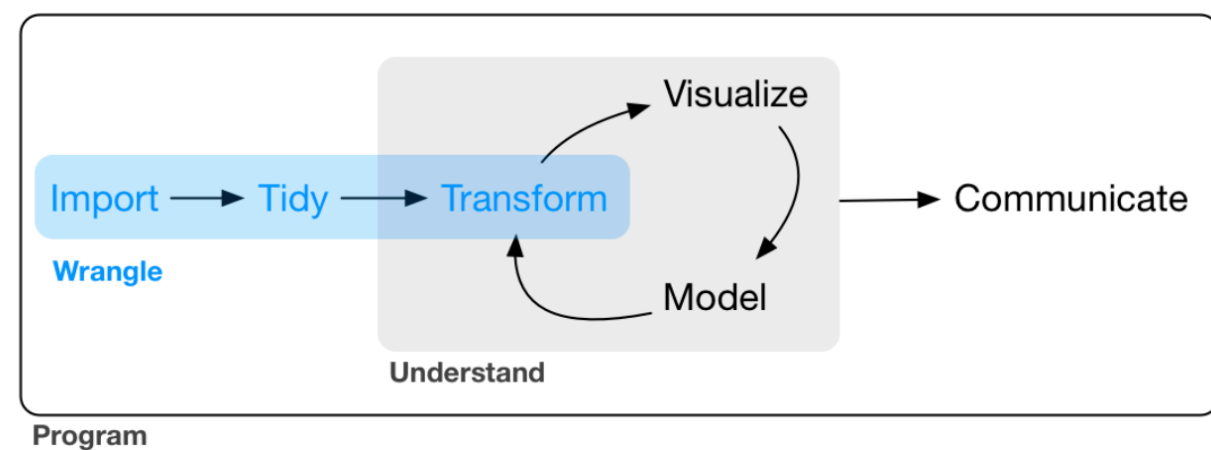
Table 1: Typical presentation dataset.

|            | John Smith | Jane Doe | Mary Johnson |
|------------|------------|----------|--------------|
| treatmenta | —          | 16       | 3            |
| treatmentb | 2          | 11       | 1            |

Table 2: The same data as in Table 1 but structured differently.

## Tidy data

1. Each variable forms a column.
2. Each observation forms a row.
3. Each type of observational unit forms a table.



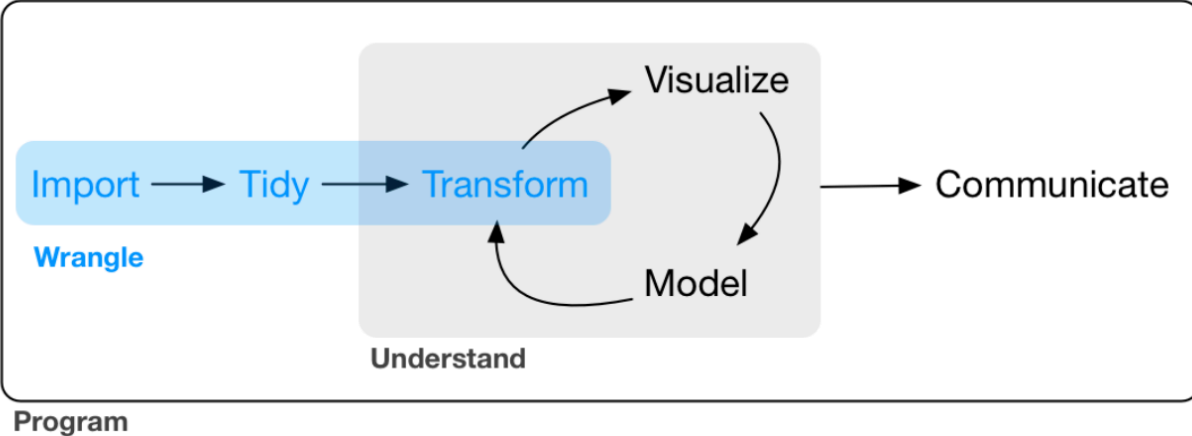
*Tidy Data*

| person       | treatment | result |
|--------------|-----------|--------|
| John Smith   | a         | —      |
| Jane Doe     | a         | 16     |
| Mary Johnson | a         | 3      |
| John Smith   | b         | 2      |
| Jane Doe     | b         | 11     |
| Mary Johnson | b         | 1      |


- Column headers are values, not variable names.
- Multiple variables are stored in one column.
- Variables are stored in both rows and columns.
- Multiple types of observational units are stored in the same table.
- A single observational unit is stored in multiple tables.

# Tidyverse

## Tidy



| religion                | <\$10k | \$10–20k | \$20–30k | \$30–40k | \$40–50k | \$50–75k |
|-------------------------|--------|----------|----------|----------|----------|----------|
| Agnostic                | 27     | 34       | 60       | 81       | 76       | 137      |
| Atheist                 | 12     | 27       | 37       | 52       | 35       | 70       |
| Buddhist                | 27     | 21       | 30       | 34       | 33       | 58       |
| Catholic                | 418    | 617      | 732      | 670      | 638      | 1116     |
| Don't know/refused      | 15     | 14       | 15       | 11       | 10       | 35       |
| Evangelical Prot        | 575    | 869      | 1064     | 982      | 881      | 1486     |
| Hindu                   | 1      | 9        | 7        | 9        | 11       | 34       |
| Historically Black Prot | 228    | 244      | 236      | 238      | 197      | 223      |
| Jehovah's Witness       | 20     | 27       | 24       | 24       | 21       | 30       |
| Jewish                  | 19     | 19       | 25       | 25       | 30       | 95       |



| religion | income             | freq |
|----------|--------------------|------|
| Agnostic | <\$10k             | 27   |
| Agnostic | \$10–20k           | 34   |
| Agnostic | \$20–30k           | 60   |
| Agnostic | \$30–40k           | 81   |
| Agnostic | \$40–50k           | 76   |
| Agnostic | \$50–75k           | 137  |
| Agnostic | \$75–100k          | 122  |
| Agnostic | \$100–150k         | 109  |
| Agnostic | >150k              | 84   |
| Agnostic | Don't know/refused | 96   |

| id      | year | month | element | d1 | d2   | d3   | d4 | d5   | d6 | d7 | d8 |
|---------|------|-------|---------|----|------|------|----|------|----|----|----|
| MX17004 | 2010 | 1     | tmax    | —  | —    | —    | —  | —    | —  | —  | —  |
| MX17004 | 2010 | 1     | tmin    | —  | —    | —    | —  | —    | —  | —  | —  |
| MX17004 | 2010 | 2     | tmax    | —  | 27.3 | 24.1 | —  | —    | —  | —  | —  |
| MX17004 | 2010 | 2     | tmin    | —  | 14.4 | 14.4 | —  | —    | —  | —  | —  |
| MX17004 | 2010 | 3     | tmax    | —  | —    | —    | —  | 32.1 | —  | —  | —  |
| MX17004 | 2010 | 3     | tmin    | —  | —    | —    | —  | 14.2 | —  | —  | —  |
| MX17004 | 2010 | 4     | tmax    | —  | —    | —    | —  | —    | —  | —  | —  |
| MX17004 | 2010 | 4     | tmin    | —  | —    | —    | —  | —    | —  | —  | —  |
| MX17004 | 2010 | 5     | tmax    | —  | —    | —    | —  | —    | —  | —  | —  |
| MX17004 | 2010 | 5     | tmin    | —  | —    | —    | —  | —    | —  | —  | —  |

| id      | date       | tmax | tmin |
|---------|------------|------|------|
| MX17004 | 2010-01-30 | 27.8 | 14.5 |
| MX17004 | 2010-02-02 | 27.3 | 14.4 |
| MX17004 | 2010-02-03 | 24.1 | 14.4 |
| MX17004 | 2010-02-11 | 29.7 | 13.4 |
| MX17004 | 2010-02-23 | 29.9 | 10.7 |
| MX17004 | 2010-03-05 | 32.1 | 14.2 |
| MX17004 | 2010-03-10 | 34.5 | 16.8 |
| MX17004 | 2010-03-16 | 31.1 | 17.6 |
| MX17004 | 2010-04-27 | 36.3 | 16.7 |
| MX17004 | 2010-05-27 | 33.2 | 18.2 |

# Tidy



variables

observations

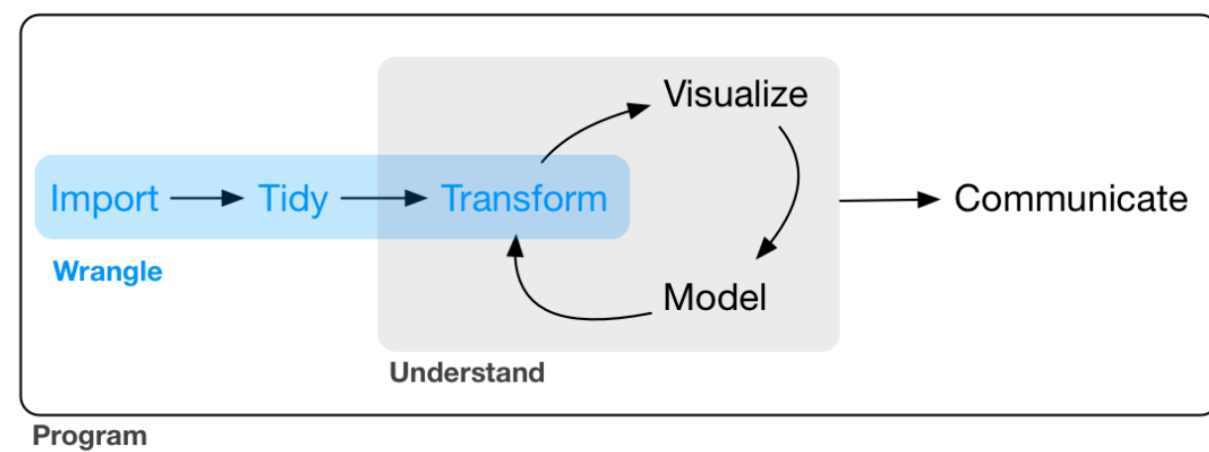
values

- Tidy dataset:
  - Variables are in columns
  - Observations in rows
  - Values in cells



# Tidyverse

## Tidy



- Spreading and Gathering
- Dataset where columns are values of a variable

```
```{r tidy1}
table4a
```
```

A tibble: 3 × 3

| country<br><chr> | 1999<br><dbl> | 2000<br><dbl> |
|------------------|---------------|---------------|
| Afghanistan      | 745           | 2666          |
| Brazil           | 37737         | 80488         |
| China            | 212258        | 213766        |

3 rows

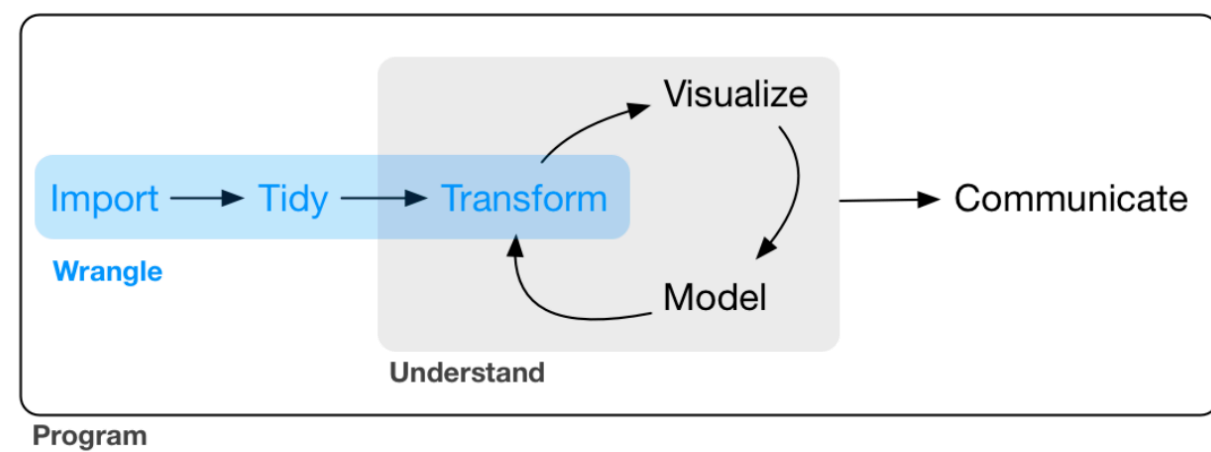
```
```{r tidy1}
table4a %>% gather(`1999`, `2000`, key="year", value="population")
```
```

A tibble: 6 × 3

| country<br><chr> | year<br><chr> | population<br><dbl> |
|------------------|---------------|---------------------|
| Afghanistan      | 1999          | 745                 |
| Brazil           | 1999          | 37737               |
| China            | 1999          | 212258              |
| Afghanistan      | 2000          | 2666                |
| Brazil           | 2000          | 80488               |
| China            | 2000          | 213766              |

# Tidyverse

## Tidy



- Spreading
- Opposite of gathering

| country     | year | key        | value      |
|-------------|------|------------|------------|
| Afghanistan | 1999 | cases      | 745        |
| Afghanistan | 1999 | population | 19987071   |
| Afghanistan | 2000 | cases      | 2666       |
| Afghanistan | 2000 | population | 20595360   |
| Brazil      | 1999 | cases      | 37737      |
| Brazil      | 1999 | population | 172006362  |
| Brazil      | 2000 | cases      | 80488      |
| Brazil      | 2000 | population | 174504898  |
| China       | 1999 | cases      | 212258     |
| China       | 1999 | population | 1272915272 |
| China       | 2000 | cases      | 213766     |
| China       | 2000 | population | 1280428583 |

table2

```

spreading
tidy2 %>% spread(key=type, value=count)

```

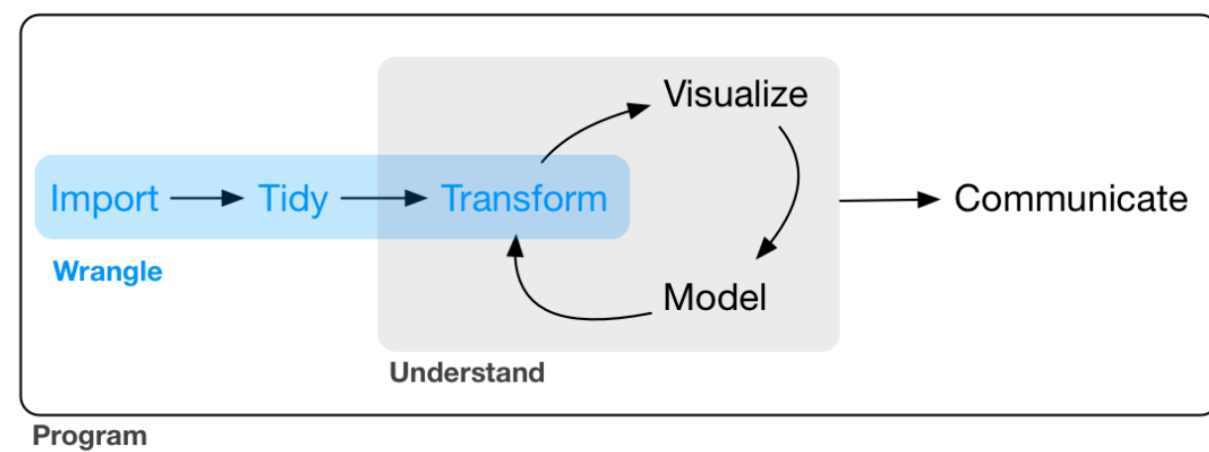
A tibble: 6 × 4

| country<br><chr> | year<br><dbl> | cases<br><dbl> | population<br><dbl> |
|------------------|---------------|----------------|---------------------|
| Afghanistan      | 1999          | 745            | 19987071            |
| Afghanistan      | 2000          | 2666           | 20595360            |
| Brazil           | 1999          | 37737          | 172006362           |
| Brazil           | 2000          | 80488          | 174504898           |
| China            | 1999          | 212258         | 1272915272          |
| China            | 2000          | 213766         | 1280428583          |



# Tidyverse

## Tidy



- `Separate()` and `unite()`

| country     | year | rate                |
|-------------|------|---------------------|
| Afghanistan | 1999 | 745 / 19987071      |
| Afghanistan | 2000 | 2666 / 20595360     |
| Brazil      | 1999 | 37737 / 172006362   |
| Brazil      | 2000 | 80488 / 174504898   |
| China       | 1999 | 212258 / 1272915272 |
| China       | 2000 | 213766 / 1280428583 |

table3

| country     | year | cases  | population |
|-------------|------|--------|------------|
| Afghanistan | 1999 | 745    | 19987071   |
| Afghanistan | 2000 | 2666   | 20595360   |
| Brazil      | 1999 | 37737  | 172006362  |
| Brazil      | 2000 | 80488  | 174504898  |
| China       | 1999 | 212258 | 1272915272 |
| China       | 2000 | 213766 | 1280428583 |

| country     | year | rate                |
|-------------|------|---------------------|
| Afghanistan | 1999 | 745 / 19987071      |
| Afghanistan | 2000 | 2666 / 20595360     |
| Brazil      | 1999 | 37737 / 172006362   |
| Brazil      | 2000 | 80488 / 174504898   |
| China       | 1999 | 212258 / 1272915272 |
| China       | 2000 | 213766 / 1280428583 |

| country     | century | year | rate                |
|-------------|---------|------|---------------------|
| Afghanistan | 19      | 99   | 745 / 19987071      |
| Afghanistan | 20      | 0    | 2666 / 20595360     |
| Brazil      | 19      | 99   | 37737 / 172006362   |
| Brazil      | 20      | 0    | 80488 / 174504898   |
| China       | 19      | 99   | 212258 / 1272915272 |
| China       | 20      | 0    | 213766 / 1280428583 |

table6

**Questions?**  
**Practice!!!**